

# Basic Population Genetics Analyses in R

Laboratory Exercise: Distributed Graduate Seminar

Rodney J. Dyer

## Overview

In this lab we will be working with some genetic marker data to:

1. Learn about basic genetic data types in R as implemented in the `gstudio` package.
2. Compute summary statistics related to diversity for complete and subsets of a data set.
3. Estimate genetic distances for a set of populations and examine hypotheses of isolation by distance.
4. Measure genetic structure among populations in a nested model.

By the end of this exercise you should be able to effectively load in, manipulate, and analyze genetic data in R. This laboratory exercise will be using the R libraries `gstudio`, `ecodist`, and `pegas`, which is freely available from <http://cran.r-project.org> to download and install OR from within R you can just type:

```
> install.packages( c("gstudio", "ecodist", "pegas") )
```

Now that you have the library on your machine, when you want to use it, you must type<sup>1</sup>:

```
> require(gstudio)
```

You will need to do this only once per session.

## Genetic Data & R

### The Locus Object

The reason for the `gstudio` library is that it defines some basic data types (like numeric, logical, etc) that describe genetic marker data. At the most fundamental level is the `Locus` object. A locus is collection of alleles but can be examined as a single entity like a floating point number or a factor, etc.

```
> loc1 <- Locus( c(120,122) )
```

```
> loc1
```

```
120:122
```

```
> loc2 <- Locus( c("A", "T") )
```

```
> loc2
```

```
A:T
```

Note, that internally the alleles are translated into `character` objects. In all the functions dealing with alleles both `integer` and `character` arguments are accepted. There are several methods associated with the `Locus`, the main ones that you will be working with are shown below by example. See `help("Locus-class")` for a complete discussion.

Once we have a locus object in R, we can inquire about its genetic characteristics and access individual alleles if necessary (via the normal vector square bracket notation `[]` just like it was a regular vector.)

```
> loc3 <- Locus( c(122,122) )
```

```
> loc3
```

```
122:122
```

```
> is.heterozygote( loc3 )
```

---

<sup>1</sup>R has a *huge* number of libraries available but allows you to only keep in memory those that you need to use.

```
[1] FALSE
> loc3[2]
[1] "122"
> loc3[2] <- "124"
> is.heterozygote( loc3 )
[1] TRUE
> length( loc3 )
[1] 2
> summary( loc3 )
Class : Locus
Ploidy : 2
Alleles : 122,124
```

## The Population Object

Having a single locus is dandy, but where we want to be is to be working with populations of individuals, each of which has the potential to have several loci. In R, the `data.frame` is a general component that you keep your data in for analyses. The `gstudio` packages defines a `data.frame` for loci called a `Population`. A `Population` object can contain all the kinds of data that we typically have, such as:

1. Stata (population names, regions, habitats, etc)
2. Spatial coordinates (latitude & longitude)
3. Other covariates (slope, soil moisture, time since colonization, etc.)
4. Genotypes (haploid, diploid, dominant, etc.)

In this next example, I will define a `Population` object consisting of a set of populations, and environmental variable measured at each individual, and two loci.

```
> strata <- c("Cabo", "Cabo", "Loreto", "Loreto", "Loreto")
> TPI <- c(Locus(c(1,2)), Locus(c(2,3)), Locus(c(2,2)), Locus(c(2,2)), Locus(c(1,3)))
> PGM <- c(Locus(c(4,4)), Locus(c(4,3)), Locus(c(4,4)), Locus(c(3,4)), Locus(c(3,3)))
> Env <- c(12,20,14,18,10)
> thePop <- Population( Pop=strata, Env=Env, TPI=TPI, PGM=PGM )
> thePop
```

```
      Pop Env TPI PGM
1  Cabo  12 1:2 4:4
2  Cabo  20 2:3 3:4
3 Loreto  14 2:2 4:4
4 Loreto  18 2:2 3:4
5 Loreto  10 1:3 3:3
```

```
> summary(thePop)
```

Pop	Env	TPI	PGM
Length:5	Min. :10.0	1:2:1	3:3:1
Class :character	1st Qu.:12.0	1:3:1	3:4:2
Mode :character	Median :14.0	2:2:2	4:4:2
	Mean :14.8	2:3:1	
	3rd Qu.:18.0		
	Max. :20.0		

You can access elements of the population by either indexes or named columns and indexes:

```
> thePop[1,2]
[1] 12
```

```
> thePop$Pop[1:4]
[1] "Cabo"    "Cabo"    "Loreto"  "Loreto"
> thePop[ thePop$Env>12,]
      Pop Env TPI PGM
1   Cabo  20 2:3 3:4
2 Loreto  14 2:2 4:4
3 Loreto  18 2:2 3:4
```

So when you take a partition of a Population object, it creates a new object with the subset of the data you used. Often we take partitions of data based upon strata (pop, region, etc) and for this, there is an easy function, `partition()` that will return a list object whose keys are the strata levels you request. For example, using the data above, I may want to partition the whole data set into populations and then estimate parameters (heterozygosity, allele frequencies, etc) on each stratum.

```
> pops <- partition(thePop, stratum="Pop")
> class(pops)
[1] "list"
> names(pops)
[1] "Cabo"    "Loreto"
> pops[[1]]
      Env TPI PGM
1   12  1:2 4:4
2   20  2:3 3:4
> pops["Loreto"]
$Loreto
      Env TPI PGM
1   14  2:2 4:4
2   18  2:2 3:4
3   10  1:3 3:3
> class( pops$Cabo )
[1] "Population"
attr(,"package")
[1] "gstudio"
```

Notice that you can access the Population objects in the list by either index number or population name. Having a list of Population objects is very helpful in R as you can extract lots of information from the list using the `lapply()` (list apply) function. For example, if you wanted to know the size of the population data sets across all strata you could type:

```
> lapply( pops, dim )
$Cabo
[1] 2 3

$Loreto
[1] 3 3
```

Where the `dim` is a function that returns the dimension (rows & columns) of the object. Here you can see that the Cabo stratum has 2 rows and 3 cols whereas the Loreto population has 3 rows and 2 columns. The `lapply` function can take more complicated functions than things like `dim`. In the following example, I show how to get the average of the `Env` variable by population using what is called an anonymous function. Here I just make up a function that every element of the `pops` list will be passed.

```
> lapply( pops, function(x) mean( x$Env ) )
```

```
$Cabo
[1] 16
```

```
$Loreto
[1] 14
```

When R reads this, it says, hey I've got a list and I'm going to call every element of the list 'x' then I'm going to give it to function(x) and whatever it returns is what I'll print out (OK, R doesn't use subjective pronouns but you get the idea).

## Reading Data From A File

R can read a wide variety of data formats. For this exercise, we will use a dataset that is included with the `gstudio` library itself so we will not have to worry too much about importing. What follows is a general overview, of how we can get data into R from a text file. In general, if you keep your data in a spreadsheet, you will export your spreadsheet as a CSV file and import that. For more info on how to do this, there is a much longer description of these methods in the vignette. Vignettes are short PDF documents that come with R packages that go into more depth about the functionality of the library. To see all the vignettes installed, at the R prompt type `browseVignettes()` and your browser will open and give you a list.

For this exercise, we will be using data that comes with `gstudio`. This data is from the bark beetle *Araptus attenuatus* that the Dyer laboratory has been working on in Baja California. To load and examine this data set type:

```
> data( araptus_attenuatus )
> summary( araptus_attenuatus )
```

Species	Cluster	Pop	Individual	Lat
CladeA: 75	CBP-C :150	32 : 19	101_10A: 1	Min. :23.08
CladeB: 36	NBP-C : 84	75 : 11	101_1A : 1	1st Qu.:24.59
CladeC:252	SBP-C : 18	Const : 11	101_2A : 1	Median :26.25
	SCBP-A: 75	12 : 10	101_3A : 1	Mean :26.25
	SON-B : 36	153 : 10	101_4A : 1	3rd Qu.:27.53
		157 : 10	101_5A : 1	Max. :29.33
		(Other):292	(Other):357	
Long	LTRS	WNT	EN	EF
Min. :-114.3	01:01:147	03:03 :108	01:01 :225	01:01:219
1st Qu.: -113.0	01:02: 86	01:01 : 82	01:02 : 52	01:02: 52
Median :-111.5	02:02:130	01:03 : 77	02:02 : 38	02:02: 90
Mean :-111.7		02:02 : 62	03:03 : 22	NA : 2
3rd Qu.: -110.5		NA : 11	01:03 : 7	
Max. :-109.1		03:04 : 8	03:04 : 6	
		(Other): 15	(Other): 13	
ZMP	AML	ATPS	MP20	
01:01: 46	08:08 : 51	05:05 :155	05:07 : 64	
01:02: 51	07:07 : 42	03:03 : 69	07:07 : 53	
02:02:233	07:08 : 42	09:09 : 66	18:18 : 52	
NA : 33	04:04 : 41	02:02 : 30	05:05 : 48	
	NA : 23	07:09 : 14	05:06 : 22	
	07:09 : 22	08:08 : 9	11:11 : 12	
	(Other):142	(Other): 20	(Other):112	

You can see that these data have categorical strata (species, cluster, pop, individual), spatial data (lat & long), and 8 codominant loci.

For the rest of the text, I am just going to use a subset of the data and allow you to answer some questions using the larger data set at the end. The subset I am going to use is the genotypes and populations from "Clade A" & "Clade B" (e.g., those in Baja California) and I'm going to call it 'pops' for brevity.

```
> pops <- araptus_attenuatus[ araptus_attenuatus$Species != "CladeB", ]
> summary(pops)
```

Species	Cluster	Pop	Individual	Lat
CladeA: 75	CBP-C :150	75 : 11	12_10A : 1	Min. :23.08

CladeC:252	NBP-C : 84	Const : 11	12_1A : 1	1st Qu.:24.59
	SBP-C : 18	12 : 10	12_2A : 1	Median :26.02
	SCBP-A: 75	153 : 10	12_3A : 1	Mean :26.18
		157 : 10	12_4A : 1	3rd Qu.:27.53
		160 : 10	12_5A : 1	Max. :29.33
		(Other):265	(Other):321	

Long	LTRS	WNT	EN	EF
Min. :-114.3	01:01:146	03:03 :108	01:01 :218	01:01:196
1st Qu.: -113.0	01:02: 69	01:03 : 76	01:02 : 52	01:02: 41
Median :-111.7	02:02:112	02:02 : 62	02:02 : 38	02:02: 90
Mean :-111.9		01:01 : 53	01:03 : 5	
3rd Qu.: -110.7		03:04 : 8	01:04 : 4	
Max. :-109.6		01:04 : 7	03:03 : 3	
		(Other): 13	(Other): 7	

ZMP	AML	ATPS	MP20
01:01: 45	08:08 : 50	05:05 :155	05:07 :64
01:02: 51	07:07 : 42	03:03 : 69	07:07 :53
02:02:214	07:08 : 42	09:09 : 65	18:18 :52
NA : 17	04:04 : 41	07:09 : 14	05:05 :48
	07:09 : 22	08:08 : 9	05:06 :22
	08:09 : 22	03:06 : 3	06:06 :11
	(Other):108	(Other): 12	(Other):77

## Allele Frequencies

One of the first things we do when examining population genetic structure is to look at allele frequencies. The `gstudio` package defines the `AlleleFrequency` data type that allows us to get individual frequencies, heterozygosities, etc.

You can get allele frequencies from a single locus

```
> freqs <- allele.frequencies( pops, loci="AML")
> freqs
```

```
$AML
```

```
Allele Frequencies:
```

```
08 = 0.2664577
09 = 0.1551724
07 = 0.2507837
10 = 0.01410658
06 = 0.07680251
11 = 0.001567398
02 = 0.001567398
13 = 0.001567398
05 = 0.07523511
01 = 0.001567398
03 = 0.02037618
04 = 0.1347962
```

a subset of loci

```
> freqs <- allele.frequencies( pops, loci=c("ATPS","LTRS"))
> freqs
```

```
$ATPS
```

```
Allele Frequencies:
```

```
09 = 0.2262997
07 = 0.03058104
05 = 0.4816514
10 = 0.004587156
03 = 0.2155963
02 = 0.006116208
```

```
08 = 0.02752294
01 = 0.003058104
06 = 0.004587156
```

```
$LTRS
```

```
Allele Frequencies:
```

```
01 = 0.5519878
02 = 0.4480122
```

```
or from all loci
```

```
> freqs <- allele.frequencies( pops )
```

```
> names(freqs)
```

```
[1] "LTRS" "WNT" "EN" "EF" "ZMP" "AML" "ATPS" "MP20"
```

```
> freqs$MP20
```

```
Allele Frequencies:
```

```
07 = 0.2892308
05 = 0.2969231
15 = 0.001538462
08 = 0.02769231
06 = 0.08923077
04 = 0.009230769
18 = 0.1784615
19 = 0.009230769
17 = 0.04153846
10 = 0.01384615
11 = 0.04153846
16 = 0.001538462
```

Once you have a frequency object, you can get both observed and expected heterozygosity (and inbreeding  $F$  if you like).

```
> ho( freqs$MP20 )
```

```
ho
```

```
0.4246154
```

```
> he( freqs$MP20 )
```

```
he
```

```
0.783787
```

```
> f <- 1- ho( freqs$MP20 ) / he( freqs$MP20 )
```

```
> f
```

```
ho
```

```
0.4582515
```

**Warning:** the `allele.frequencies()` function returns a list indexed by the name of the locus even if you only are asking for a single locus. We will see this behavior again when we estimate genetic distances.

So, now, it is possible to do something fun with the data. Lets plot the location of the populations again but this time lets make the symbol size scaled by  $H_e$ . This is the entire species' range we are examining and there are several kinds of core vs. periphery hypotheses that can be examined. What we will do is first use `lapply` as above to get  $H_e$  for all AML loci by population (and then take it out of a list using the `unlist`) function). Options I pass to the plot are `pch` which is the plot character (a filled circle), `cex` is the character expansion size (scaling of the symbol), and `bty` is the box type around the plot (I personally hate boxes around my plots). I scale the  $H_e$  values for visual differences because the size is proportional to heterozygosity, not the latitude/longitude.

```
> subpops <- partition( pops, stratum="Pop")
```

```
> he <- unlist(lapply( subpops, function(x) he(allele.frequencies(x,"AML")[[1]] ) ))
```


```
> he
```

```

12.he 153.he 156.he 157.he 159.he 160.he 161.he 162.he
0.620000 0.515000 0.444444 0.460000 0.5679012 0.725000 0.635000 0.585000
163.he 164.he 165.he 166.he 168.he 169.he 171.he 173.he
0.460000 0.485000 0.625000 0.790000 0.615000 0.605000 0.595000 0.635000
175.he 177.he 48.he 51.he 58.he 64.he 73.he 75.he
0.5816327 0.675000 0.180000 0.500000 0.611111 0.540000 0.640000 0.593750
77.he 84.he 88.he 89.he 9.he 93.he 98.he Aqu.he
0.535000 0.4609375 0.6728395 0.650000 0.444444 0.610000 0.375000 0.6428571
Const.he ESan.he Mat.he SFr.he
0.685000 0.531250 0.500000 0.6851852

> scaled_he <- 2*he + .5
> lat <- unique(pops$Lat);
> lon <- unique(pops$Long)
> plot( lat, lon, xlab="Latitude", ylab="Longitude", pch=16, cex=scaled_he, bty="n")

```



## Genetic Diversity

There are several measures of genetic diversity, many of which we can quickly estimate in R.

**Polymorphic Loci:** The number of polymorphic loci. We can go through each population in the data set and for each locus determine if there are more than one allele (tedious). Or we can make R do it with a little looping or cleaverness (much more exciting). First I am going to define a function that takes a list of `AlleleFrequency` objects (what you get from a call to `allele.frequencies`) and finds out how many loci have a length  $> 1$  (e.g., they have more

than one allele). It returns the value as a fraction of how many total loci there are.

```
> polymorphic.loci <- function(x) {
+   sum(lapply( x, length ) > 1) / sum(length(x))
+ }
```

Next, I can apply this to all the data.

```
> all.freqs <- lapply( subpops, allele.frequencies )
> P <- lapply( all.freqs, polymorphic.loci )
> unlist(P)
```

	12	153	156	157	159	160	161	162	163	164	165	166	168
0.750	0.625	0.625	0.625	0.875	0.875	0.625	1.000	0.750	1.000	0.875	1.000	0.750	
	169	171	173	175	177	48	51	58	64	73	75	77	84
0.750	0.875	0.875	0.625	0.625	0.375	0.625	0.625	0.500	0.625	0.875	0.750	0.875	
	88	89	9	93	98	Aqu	Const	ESan	Mat	SFr			
0.750	0.750	0.625	0.750	0.625	0.875	1.000	0.625	0.750	0.750				

This shows one of the real strengths of R. What we've done here is made a completely new function and can use it as long as it is in memory or can save it to our own personal library of functions.

**Allelic Diversity:** Allelic diversity can be quantified as the number of alleles at a locus ( $A$ ), the number of alleles at a locus above some predefined frequency ( $A_{95}$  for those whose frequencies exceed 5%), and the effective number of alleles ( $A_e$ ). The `genetic.diversity` function allows you to estimate all of these parameters and does so using rarefaction (e.g., permutations to test diversity for standardized sample sizes, it may take a minute to do it on your computer).

```
> diversity.mp20 <- genetic.diversity( pops, stratum="Pop", mode="Ae", loci="MP20")
> diversity.mp20
```

Geneic Diversity:

```
Estimator: Ae
Stratum: Pop
Loci: { MP20 }
Locus = MP20
12 Ae = 2.98507462686567 ; Rarefaction Ae = 2.576352595699
153 Ae = 1.69491525423729 ; Rarefaction Ae = 1.70087039077751
156 Ae = 2.18181818181818 ; Rarefaction Ae = 2.02766271348315
157 Ae = 2.73972602739726 ; Rarefaction Ae = 2.34135280314847
159 Ae = 2.41791044776119 ; Rarefaction Ae = 2.11275046969577
160 Ae = 3.33333333333333 ; Rarefaction Ae = 2.78634896447074
161 Ae = 4.16666666666667 ; Rarefaction Ae = 3.13368481814302
162 Ae = 3.2258064516129 ; Rarefaction Ae = 2.62689419958082
163 Ae = 2.17391304347826 ; Rarefaction Ae = 2.01553350560574
164 Ae = 1.85185185185185 ; Rarefaction Ae = 1.78996445997478
165 Ae = 3.17460317460317 ; Rarefaction Ae = 2.59848584057734
166 Ae = 3.125 ; Rarefaction Ae = 2.72344186794823
168 Ae = 2.89855072463768 ; Rarefaction Ae = 2.45144682032195
169 Ae = 2.1978021978022 ; Rarefaction Ae = 2.0421096990206
171 Ae = 2.1978021978022 ; Rarefaction Ae = 1.96317772751209
173 Ae = 2.46913580246914 ; Rarefaction Ae = 2.19183418199586
175 Ae = 1.96 ; Rarefaction Ae = 1.79530118353648
177 Ae = 1.8018018018018 ; Rarefaction Ae = 1.72853823491594
48 Ae = 1.6 ; Rarefaction Ae = 1.55902726255667
51 Ae = 1.55555555555556 ; Rarefaction Ae = 1.56073680527757
58 Ae = 1.90588235294118 ; Rarefaction Ae = 1.83654635084976
64 Ae = 1 ; Rarefaction Ae = 1
73 Ae = 2.17391304347826 ; Rarefaction Ae = 2.07897901095011
75 Ae = 2.3047619047619 ; Rarefaction Ae = 2.1009077435634
77 Ae = 2.10526315789474 ; Rarefaction Ae = 2.03445786303715
```



```

84 Ae = 1.90588235294118 ; Rarefaction Ae = 1.80051656121832
88 Ae = 1.6 ; Rarefaction Ae = 1.55172348819408
89 Ae = 2.06185567010309 ; Rarefaction Ae = 1.89543970175549
9 Ae = 1.11724137931034 ; Rarefaction Ae = 1.11980686568922
93 Ae = 2.98507462686567 ; Rarefaction Ae = 2.54567897347395
98 Ae = 1.68421052631579 ; Rarefaction Ae = 1.65878854006821
Aqu Ae = 2.8 ; Rarefaction Ae = 2.4462072725437
Const Ae = 2.37254901960784 ; Rarefaction Ae = 2.1865324214963
ESan Ae = 2.24561403508772 ; Rarefaction Ae = 2.12420737288986
Mat Ae = 1.28 ; Rarefaction Ae = 1.29617853147265
SFr Ae = 2.18918918918919 ; Rarefaction Ae = 2.00934647747169

```

**Heterozygosities:** In the previous section we went into heterozygosity estimation by populations so we will probably not need to cover it again here.

## Genetic Distance

There are many different kinds of genetic distance and `gstudio` hides them within the function `genetic.distance`. When estimating distance, you must pass a population, the name of the stratum on which to partition, and the mode of the distance calculation. See `?genetic.distance` for more information. At present the following distance types are available:

**Individual Distances:** These are distances measured among individuals. The resulting matrices will be  $N \times N$  in size.

- *Jaccard*
- *Bray*
- *AMOVA*

**Stratum Distances:** These are estimated among strata resulting in a matrix of size  $K \times K$

- Euclidean
- Cavalli-Sforza
- Nei
- cGD

There is a corresponding distance metric that can be estimated from coordinates using `stratum.distance` that will return the "great circle distance" (or euclidean) from a set of strata. What we'll do is measure Nei's distance among populations (need a  $-1 * \log(nei)$  for standard Nei's distance) in the data set and then plot that against physical distance (IBD). When we use this, we'll test for a correlation using a Mantel test from the `ecodist` library.

```

> phys <- stratum.distance(pops, stratum="Pop", lat="Lat", lon="Long")
> nei <- genetic.distance( pops, stratum="Pop", mode="Nei")[[1]]
> require(ecodist)
> mantel( as.dist(phys) ~ as.dist(nei) )

      mantelr      pval1      pval2      pval3  llim.2.5% ulim.97.5%
-0.6232364  1.0000000  0.0010000  0.0010000 -0.6611961 -0.5846042

```

For some distance metrics, there are alternative ways to accumulate distances across loci. As such, I have left the individual loci separate and let you decide how to combine them. Here is an example of this using a `for`-loop adding all the single locus values.

```

> cav <- genetic.distance(pops, stratum="Pop", mode="Cavalli")
> names(cav)

[1] "LTRS" "WNT" "EN" "EF" "ZMP" "AML" "ATPS" "MP20"

> cav.all.loci <- matrix(0,nrow=36,ncol=36)
> for( locus in names(cav) )
+     cav.all.loci <- cav.all.loci + cav[[locus]]
> mantel( as.dist(cav.all.loci) ~ as.dist(phys) )

```

```

mantelr      pval1      pval2      pval3  llim.2.5% ulim.97.5%
0.6333091  0.0010000  1.0000000  0.0010000  0.5905275  0.7010702

```

## Genetic Structure

Genetic structure is a measure of among-strata configuration. In the lecture we examined  $F_{ST}$ ,  $G_{ST}$ , &  $\Theta$  as population parameters,  $G'_{ST}$  &  $D_{est}$  as population parameters standardized for highly diverse loci, and  $\Phi_{ST}$  as a multilocus statistical measure of differentiation. Population structure parameters are fundamental tools for population genetics and have been perhaps, the most poorly understood and misused as well.

These structure parameters are estimated using the function `genetic.structure` and requires a `Population` object, a `stratum`, the loci you want to estimate parameters from, and a `mode` (the parameter you want). If you leave off the `loci` parameter, all loci will be used. There is also an optional parameter, `num.perm` that is used to test significance. In what follows, we will examine the following parameters as a demonstration of how to estimate these parameters:

$G_{ST}$ : This parameter is estimated from the differences in observed and expected heterozygosity.

```

> gst <- genetic.structure( pops, stratum="Pop", loci="EN", mode="Gst", num.perm=0)
> gst

```

Geneic Structure Analysis:

```

Estimator: Gst
Stratum: Pop
Loci: { EN }
- EN ; Gst = 0.345786963051191

```

$G'_{ST}$  &  $D_{Est}$ : These parameters provide corrections that are caused by loci with high allelic diversity.

```

> gst.prime <- genetic.structure(pops, stratum="Pop", loci="EN", mode="Gst.prime", num.perm=0 )
> gst.prime

```

Geneic Structure Analysis:

```

Estimator: Gst.prime
Stratum: Pop
Loci: { EN }
- EN ; Gst.prime = 0.459618108931204

```

```

> dest <- genetic.structure(pops, stratum="Pop",loci="EN",mode="Dest", num.perm=0)
> dest

```

Geneic Structure Analysis:

```

Estimator: Dest
Stratum: Pop
Loci: { EN }
- EN ; Dest = 0.164464814867075

```

## Pairwise Structure

There are times when we may be intersted in estimating pairwise structure parameters. This can be done by invoding the optional pairwise flag for `genetic.structure`. Here is an example using the populations in Sonora (e.g., the non-peninsular populations).

```

> sonora <- araptus_attenuatus[araptus_attenuatus$Species=="CladeB",]
> genetic.structure(sonora,"Pop",loci="EN",mode="Gst.prime", pairwise=TRUE)

```

```

          101          32          102
101 0.0000000 0.4136727 0.3661894
32  0.4136727 0.0000000 0.1245850
102 0.3661894 0.1245850 0.0000000

```

## AMOVA

The AMOVA is a statistical decomposition of genetic variance into two additive components, the within stratum component ( $\sigma_w^2$ ) and the among strata component ( $\sigma_A^2$ ). AMOVA produces a statistic,  $\Phi_{ST}$  which is defined as:

$$\Phi_{ST} = \frac{\sigma_A^2}{\sigma_A^2 + \sigma_W^2}$$

using a distance matrix approach. Here we will use the smaller sonora data set to decompose genetic structure and estimate this parameter.

```
> require(pegas)
> D <- genetic.distance( sonora, mode="AMOVA")[[1]]
> D <- as.dist(D)
> Pops <- as.factor( sonora$Pop )
> fit.amova <- amova(D ~ Pops)
> summary(fit.amova)
```

	Length	Class	Mode
tab	3	data.frame	list
varcoef	1	-none-	numeric
varcomp	2	data.frame	list
call	2	-none-	call

```
> fit.amova
```

Analysis of Molecular Variance

Call: amova(formula = D ~ Pops)

	SSD	MSD	df
Pops	445.7229	222.86146	2
Error	919.2489	27.85603	33
Total	1364.9719	38.99920	35

Variance components:

	sigma2	P.value
Pops	17.773	0
Error	27.856	

Variance coefficients:

a  
10.97222

Now unfortunately, the parameter  $\Phi_{ST}$  is not directly produced by the output (don't know why but that is the way the author of the pegas library wrote it). However, it is easily calculated as:

```
> sigmaA <- fit.amova$varcomp[1,1]
> sigmaW <- fit.amova$varcomp[2,1]
> Phi <- sigmaA / (sigmaA+sigmaW)
> Phi
```

```
[1] 0.3895061
```

## Exercises

The following exercises will allow you to test out the basic skills you learned in this laboratory.

```
> require(gstudio)
> data(araptus_attenuatus)
> data <- araptus_attenuatus[ araptus_attenuatus$Species=="CladeC",]
> counts <- table(data$Pop)
> counts
```

12	153	157	159	160	161	162	163	164	165	166	168	169
10	10	2	9	10	10	10	7	8	10	8	10	10
171	173	175	177	51	58	64	73	75	77	84	88	89
10	10	7	10	7	9	5	2	1	9	9	10	10
9	93	98	Aqu	Const	ESan	Mat	SFr					
9	10	1	4	3	2	1	9					

If we look at this data, we can see we have a variable number of samples per population. In fact, for this Clade, there are several species with small sample sizes (as it turns out this is because what we thought was one species is actually two separate species in sympatry). So let's go through the data and remove those populations with fewer than 5 samples. If you look at the variable counts it is a numeric vector whose names are the population names. From this, we can find the population names whose counts are greater than 5

```
> keepers <- names(counts[ counts > 5 ])
> keepers
```

```
[1] "12" "153" "159" "160" "161" "162" "163" "164" "165" "166" "168" "169"
[13] "171" "173" "175" "177" "51" "58" "77" "84" "88" "89" "9" "93"
[25] "SFr"
```

And then only use the data from those populations using the %in% operator.

```
> data <- data[ data$Pop %in% keepers, ]
> table(data$Pop)
```

12	153	159	160	161	162	163	164	165	166	168	169	171	173	175	177	51	58	77	84
10	10	9	10	10	10	7	8	10	8	10	10	10	10	7	10	7	9	9	9
88	89	9	93	SFr															
10	10	9	10	9															

This is pretty cool stuff because you can easily envision how easy it is to work with various subsets of your data set. OK, now on to some questions.

1. Can you rank these populations in terms of genetic diversity? What metric did you choose and why?
2. In the Clade C data, is there any indication of changes in expected heterozygosity as a function of either latitude or longitude? You can use the `cor.test()` function to test for significance.
3. In addition to strata-level genetic distances, there are also several individual-level genetic distance measures available. How correlated are the individual genetic distances from methods such as "AMOVA" and "Jaccard"? You may want to use the `mantel` function from the `ecodist` library as we did for population-level distances. Also, since the Jaccard distance is a single-locus estimates, you can either combine them across loci for a multilocus estimate or look at the loci individually.
4. I didn't use Bray-Curtis in the previous question because there are some missing data. Can you think of a way to handle missing data using this metric so that a comparison can be made?
5. Of the single-locus measures of genetic structure, which one would you use to estimate among-population structure and why? Is there a lot of structure in these data or a little?